


Capital Group

POV via. Ahead



Table of contents

1. Current state of data security at Capital Group
2. The S.A.F.E approach to securing data
3. S.A.F.E performance with Bedrock in Ahead PoV
4. Key out of box findings
 - a. Data discovery
 - b. Sensitive data access
 - c. Sensitive data leakage
5. Roadmap
6. Stakeholder value, present and future
7. Appendix / Requirements

 In this report, each underlined segment is a deep link into Bedrock that can be clicked.

Try it out!

Current state of data security at Capital Group

- Incomplete visibility into critical data across all data platforms

💬 *“We don't know what the data looks like. We don't know where our data is.” - Khanh*

↳ **Data policy violations, leaks or exposures may be unnoticed**

- Manual data classification and tagging lacks accuracy (Resource Constraints)

💬 *“[User] might say this is non-sensitive, but then six months later it's sensitive, they don't go back and change that data tagging” - Khanh*

↳ **Unable to write accurate data policy/controls to keep data safe w/ growth and GenAI**

- Limited visibility into data access, usage, and movement (Data Management)

💬 *“What kind of sharing policies do we have, but also show me all the data that's associated with this particular identity or line of business.” - Khanh*

↳ **Cannot track exposure of data as access levels change or as data is moved**

The **S.A.F.E** approach to securing data

- Scalable** Data should be secured without gaps in multi PB environments such as CG's Caspian data lake
- Accurate** Identification of sensitive data and risks must be reliable for both regulatory data (e.g. NPI) and proprietary business data (e.g. financial trades)
- Fast** Data scanning and risk assessment techniques must deliver rapid time to value and capture exposures / control violations in near real-time
- Efficient** Data security initiatives should not result in costly infrastructure bills

S.A.F.E performance with Bedrock in Ahead PoV

Scalable

600TB scanned daily for 1 week

- AWS, Azure, Snowflake, Databricks & M365

Accurate

Regulated data and proprietary data identified

- PII/NPI, PCI, Financial Transactions Records and Invoices
- Transactions and Invoices were auto-discovered by Bedrock's ML

Fast


Under 9 hours required to scan 600TB

Efficient

Order of magnitude more cost efficient scanning

Key out of box findings: **Data discovery**

| Capital Group's Custom Data Learned by Bedrock | | | | | | | | | Regulated Data | | |
|--|------------|----------------|----------------|----------------|----------------|---------|------------|---------------|----------------|-------------|------|
| | | Secrets / Keys | Analytics Data | Developer Data | Financial Data | HR Data | Sales Data | Security Data | Logs | PII/PI/NPI† | PCI† |
| Infrastructure / Staging | AWS S3 | | | 33 | 7† | | | | 8 | 43 | 40 |
| | AWS DBs* | | | | | | | | | 3 | 3 |
| | AWS EFS | | | 1 | | | | | | 1 | 1 |
| | Azure Blob | | | | | | | | | 3 | 3 |
| | Snowflake | | | | | | | | | 1 | 1 |
| Corp | OneDrive | | 19 | 19 | 1 | 9 | 6 | 2 | | 93 | |
| | Sharepoint | 1 | 3 | | 4 | 2 | 4 | | | 227 | |

 Colors represent relative risk for data in a given location



* Databases include AWS RDS and Unmanaged SQL

† Identified data is likely synthetic / simulated

Key out of box findings: Custom data discovered

Bedrock's ML automatically categorizes and labels the data it scans. Below are examples of data custom to the Ahead PoV environment that Bedrock identified, including simulated financial transactions.

| Data category | Examples of custom data types learned by Bedrock |
|----------------|---|
| Secrets / Keys | <u>AWS Credentials</u> , <u>Database Credentials</u> |
| Financial Data | <u>Invoice Details</u> , <u>Purchase Orders</u> , <u>Tax Documents</u> , <u>Synthetic Financial Transactions</u> |
| HR Data | <u>Health Authorization</u> , <u>Resume</u> , <u>Onboarding Documents</u> , <u>Offshore Staffing</u> , <u>Termination Notices</u> |
| Sales Data | <u>Carrier Tracking</u> , <u>Dell Server Quotes</u> , <u>Production Orders</u> |

Key out of box findings: **Sensitive data access**

- Identified 402 identities w/ sensitive access to **PII/NPI, PCI and Transactions data**
 - ↳ **Decommission 2 stale, high impact human identities**, inactive for over 6 months
 - ↳ **Investigate 206 unknown identities from domain @dccplbng** for sensitive access
 - ↳ **Revoke PII shared with personal @gmail accounts: Uma and Sushma**

| | Infrastructure (e.g. AWS, Azure, Snowflake) | Corporate (e.g. OneDrive, Sharepoint) |
|-------------|---|---------------------------------------|
| Human | 17 (2 stale) | 152 (@dhruvts.com) |
| Non-human | 24 | 206 (@dccplbng) |
| Third-party | 1 (Databricks) | 2 (@gmail.com) |

Key out of box findings: **Sensitive data leakage**

- Sensitive database backups show up in 13 personal OneDrives / Sharepoints
 - ↳ **Delete** database backups
- PII residing in non-production assets (created by John Soto)
 - ↳ Quarantine PII in high PII volume assets using built-in remediations
- Data from service database i-04a823e99f856141b leaking outside of AWS
 - ↳ Delete data copied into Azure Blob Store

Roadmap

Integrations

- Q4 Wiz (integration already approved)
- Q4 Enhancements to Databricks e.g. unity catalog
- Y25 On-premise sources based on roadmap agreement

Platform

- Q3 Data catalog and Data Bill of Materials; (shipped)
- Q3 Chained entitlement visualization; (shipped 2nd generation)
- Q4 Least-privilege recommendations
- Y25 Investigation and remediation copilot, impact analysis

Beyond DSPM: Stakeholder value, present and future

- **Data context for security efforts**
 - SIEM and CNAPP alert correlation with sensitive data
 - Improved classification for DLP
 - Accelerated least privilege initiatives: reduce sensitive entitlements
- **Enterprise data management**
 - Automated data labeling and metadata lake
 - Data retention, deduplication, minimization
- **Generative AI**
 - Data bill of materials for model training
- **Intellectual property tracking**
 - Tracking specific sensitive data content and trust boundary management

Appendix

Metrics and screenshots



S.A.F.E performance details in Ahead PoV

Bedrock’s adaptive sampling scanned ~600 TB of data daily for 1 week. The **total time spent scanning was less than 9 hours** and the **infrastructure cost to Ahead was less than \$1,000**.

| Workload | Data Volume | Bedrock Performance |
|-----------------------|-------------|------------------------------|
| AWS (S3, EFS, RDS) | 554 TB | 2.23 Hours |
| Unmanaged SQL Server | 4.1 GB | 10 Minutes |
| Databricks on S3 | 250 MB | 2 Minutes |
| Snowflake | 67 GB | 3 Minutes |
| Azure Blob Store | 15 TB | 9 Hours (Factoring in retry) |
| OneDrive & Sharepoint | 10.6 GB | 8.13 Hours |

Secrets / keys in Sharepoint

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

nto

SharePoint:SharePoint/CG - Ahead

Overview

Contents

Data Exposure

Access Activity

Similar Data

Visualizations

Unresolved Findings

Search for Datasets

Export

More filters

Data classifications: Access

1 result

| DATASET NAME | DATA CATEGORY | USERS W/ ACCESS | THIRD PARTIES W/ ACCESS | ROLES W/ ACCESS | SIMILAR DATASETS | LAST ACCESS TIME | NOTES |
|--|--|-----------------|-------------------------|-----------------|------------------|------------------|--------------------------|
| Documents/Cloud Infrastructure/dspm aws creds cons.zip | Edit ★ Access | | 1 third party | | | N/A | Add Note |

Feature Flags

PII/PI/NPI/PCI in Snowflake

PRODUCTION capitalgroup

Search...

NAVIGATION

- Inventory
- Data Dashboard
- Findings
- Policies
- Data Catalog
- Reports

SETTINGS

- Users
- Integrations
- Authentication
- Data Type Manager
- Configurations

DSPM_DB.PUBLIC

LXVXKLV-MMB75057.DSPM_DB:PUBLIC

copy

OverviewContentsData ExposureAccess ActivitySimilar DataVisualizationsUnresolved Findings

Search for Datasets

Export

More filtersData classifications: Personal Identifiable Information (PII), Nonpublic Personal Information, Personal Information (PI)...Datastore type: Snowflake

3 results

| DATASET NAME | DATA FOUND | ADDED BY | PARTIES W/ ACCESS | ROLES W/ ACCESS | SIMILAR DATASETS | LAST ACCESS TIME | NOTES |
|-------------------------------------|---|----------|-------------------|-----------------|------------------|------------------|----------|
| PERSONALLY_IDENTIFIABLE_INFORMATION | <div>Payment Card Information</div> <div>Payment Card Number</div> | Bedrock | | 2 roles | 1 Dataset | N/A | Add Note |
| USER_INFORMATION_02 | <div>Expiration Date</div> | Bedrock | | 2 roles | | N/A | Add Note |
| USER_INFORMATION_01 | <div>Personal Information (PI)</div> <div>Passport Number</div> <div>Payment Card Number</div> <div>Address</div> | Bedrock | | 2 roles | 1 Dataset | N/A | Add Note |

Feature Flags

High impact human identities, inactive for >6mos

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

Feature Flags

Inventory

More filters

Entity: Principal

Time since last usage: Greater than 26 Weeks

Principal types: User

Accessible data categories: No Selection

Search 2 Principals

Export

2 results

| ACTIONS | NAME | IMPACT | ACCESSIBLE DATA CATEGORIES | IDENTITY ORIGIN | PRINCIPAL TYPE | LAST ACTIVITY |
|---------|---------------|---------|---|-----------------|----------------|-----------------------|
| | aws-avtr-user | ● ● ● ● | ★ Engineering Data ★ Synthetic Data +6 | AWS | USER | 9:01am, Feb 21st 2024 |
| | shiva.p | ● ● ● ● | ★ Synthetic Data ★ Personal Identifiable Information (PII) +3 | AWS | USER | 12:09am, Jul 7th 2023 |

Data shared externally w/ Uma's @gmail

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

Uma Shankar Reddy A

ROLE

Overview

Role Accessibility

Roles

Accessible Data

Principal Activity

Unresolved Findings

Cloud Provider Tags

Visualizations

Accessible Data

Showing the datastores that are accessible by this principal

Search 1 Datastore

More filters

Data classifications: No Selection

Account name: No Selection

1 results

| DATASTORE NAME | DATASTORE TYPE | IMPACT | DATA CATEGORY | ROLE(S) USED TO ACCESS | DATASET COUNT | ACCOUNT ENVIRONMENT | ACCESS TYPE | RELATIONSHIP TO DATA STORE | ACCESS TO DATASETS |
|---|----------------|--------|---|------------------------|---------------|---------------------|------------------|----------------------------|--------------------|
| SharePoint:SharePoint/Global Dhruv Team | SharePoint | | ★ Personal Identifiable Information (PII) | 0 Roles | 238 Datasets | Staging | Admin Write Read | - | Partial |
| <div><div>★ Personal Identifiable Information (PII)</div><div>★ Name</div><div>★ Email Address</div><div>Nonpublic Personal Information</div><div>Name</div><div>★ Personal Information (PI)</div><div>★ Name</div><div>★ Email Address</div></div> | | | | | | | | | |

Feature Flags

BEDROCK

SECURITY

Data shared externally w/ Sushma's @gmail

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

sushma BN

ROLE

Overview

Role Accessibility

Roles

Accessible Data

Principal Activity

Unresolved Findings

Cloud Provider Tags

Visualizations

Accessible Data

Showing the datastores that are accessible by this principal

Search 1 Datastore

More filters

Data classifications: No Selection

Account name: No Selection

1 results

| DATASTORE NAME | DATASTORE TYPE | IMPACT | DATA CATEGORY | ROLE(S) USED TO ACCESS | DATASET COUNT | ACCOUNT ENVIRONMENT | ACCESS TYPE | RELATIONSHIP TO DATA STORE | ACCESS TO DATASETS |
|---------------------------------------|----------------|--------|---|------------------------|---------------|---------------------|------------------|----------------------------|--------------------|
| OneDrive:OneDrive/Jyothi Swarup Reddy | M365 | | ★ Personal Identifiable Information (PII) | +2 0 Roles | 1 Dataset | Staging | Admin Write Read | - | Partial |

★ Personal Identifiable Information (PII)

★ Name

★ Email Address

Nonpublic Personal Information

Name

★ Personal Information (PI)

★ Email Address

★ Name

Feature Flags

Sensitive DB backups leaking to OneDrive/Sharepoint

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

Policies > Sensitive database backups should only stay in AWS

Sensitive database backups should only stay in AWS

ACTIVE

Edit Policy

Delete Policy

Pause Policy

Overview

Unresolved Findings

Entities Affected

Entities Affected

Showing all entities that are involved in the findings created by this policy

Search entity names

More filters

13 ENTITIES SHOWN

| ENTITY NAME | DATA OWNER | ENTITY CLASSIFICATION | ACCESSIBLE DATA CATEGORIES |
|---|------------|-----------------------|---|
| OneDrive:OneDrive/Akram Pasha | - | M365 | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| OneDrive:OneDrive/Arpitha SA | - | M365 | HR Data Personal Identifiable Information (PII) +2 |
| OneDrive:OneDrive/Ananthan | - | M365 | Security Data Personal Identifiable Information (PII) +2 |
| SharePoint:SharePoint/GKB Vision | - | SharePoint | Engineering Data Personal Identifiable Information (PII) +2 |
| SharePoint:SharePoint/Cytiva | - | SharePoint | Engineering Data Personal Identifiable Information (PII) +2 |
| SharePoint:SharePoint/Dhruv BI Quick Start Solution | - | SharePoint | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| SharePoint:SharePoint/Dixon AMS Internal | - | SharePoint | Personal Identifiable Information (PII) Nonpublic Personal Information |
| SharePoint:SharePoint/Dhruv Rebranding | - | SharePoint | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| SharePoint:SharePoint/Pacbio Internal | - | SharePoint | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| OneDrive:OneDrive/Dennis Mosley | - | M365 | Analytics Data Engineering Data +4 |
| OneDrive:OneDrive/Gagan C | - | M365 | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| OneDrive:OneDrive/GANSHEKAR RAAJ H | - | M365 | Personal Identifiable Information (PII) Nonpublic Personal Information +1 |
| OneDrive:OneDrive/Harish N | - | M365 | Analytics Data Personal Identifiable Information (PII) +2 |

Feature Flags

BEDROCK

SECURITY

Data shared externally w/ Sushma's @gmail

PRODUCTION capitalgroup

Search...

K

NAVIGATION

Inventory

Data Dashboard

Findings

Policies

Data Catalog

Reports

SETTINGS

Users

Integrations

Authentication

Data Type Manager

Configurations

sushma BN

ROLE

Overview

Role Accessibility

Roles

Accessible Data

Principal Activity

Unresolved Findings

Cloud Provider Tags

Visualizations

Accessible Data

Export

Showing the datastores that are accessible by this principal

Search 1 Datastore

More filters

Data classifications: No Selection

Account name: No Selection

1 results

| DATASTORE NAME | DATASTORE TYPE | IMPACT | DATA CATEGORY | ROLE(S) USED TO ACCESS | DATASET COUNT | ACCOUNT ENVIRONMENT | ACCESS TYPE | RELATIONSHIP TO DATA STORE | ACCESS TO DATASETS |
|---------------------------------------|----------------|--------|---|------------------------|---------------|---------------------|------------------|----------------------------|--------------------|
| OneDrive:OneDrive/Jyothi Swarup Reddy | M365 | | ★ Personal Identifiable Information (PII) | +2 0 Roles | 1 Dataset | Staging | Admin Write Read | - | Partial |

★ Personal Identifiable Information (PII)

★ Name

★ Email Address

Nonpublic Personal Information

Name

★ Personal Information (PI)

★ Email Address

★ Name

Feature Flags

Remediate sensitive data outside of production

The screenshot displays the Bedrock Security console interface. The top navigation bar is red and contains the 'PRODUCTION capitalgroup' label, a search bar, and a user profile icon. The left sidebar lists various navigation items: Inventory, Data Dashboard, Findings, Policies, Data Catalog, Reports, Users, Integrations, Authentication, Data Type Manager, and Configurations. The main content area is titled 'user-registration-data-usw-1 - Test - Sensitive Data should not reside in non-sensitive area' and includes a timestamp 'Oct 1st 2024, 5:48pm (1d ago)' and a status 'Unresolved'. Below the title, there are tabs for Overview, Impact Summary, Data Exposure, Access Activity, Compliance, and Remediation. The Remediation tab is active, showing a task titled 'Isolate S3 Bucket'. The task description states: 'Isolate the bucket using the following [Cloud Formation Template](#). Please add the administrators in the cloud formation template or do the following steps: 1. Create a JSON file (e.g., 'bucket_policy.json') containing the bucket policy. Replace 'administrator_arns' with the ARNs of the IAM users or groups you want to grant access to.' A code block contains the following JSON template:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "QuarantineBucket",
      "Effect": "Deny",
      "Principal": "*",
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::user-registration-data-usw-1",
        "arn:aws:s3:::user-registration-data-usw-1/*"
      ],
      "Condition": {
        "StringNotLike": {
          "aws:userId": [
            "administrator_arn",
            "administrator_arn2"
            // ...additional administrator ARNs, if any
          ]
        }
      }
    }
  ]
}
```

At the bottom left of the console, there is a 'Feature Flags' button.

Data leaving core service DB into other services

